



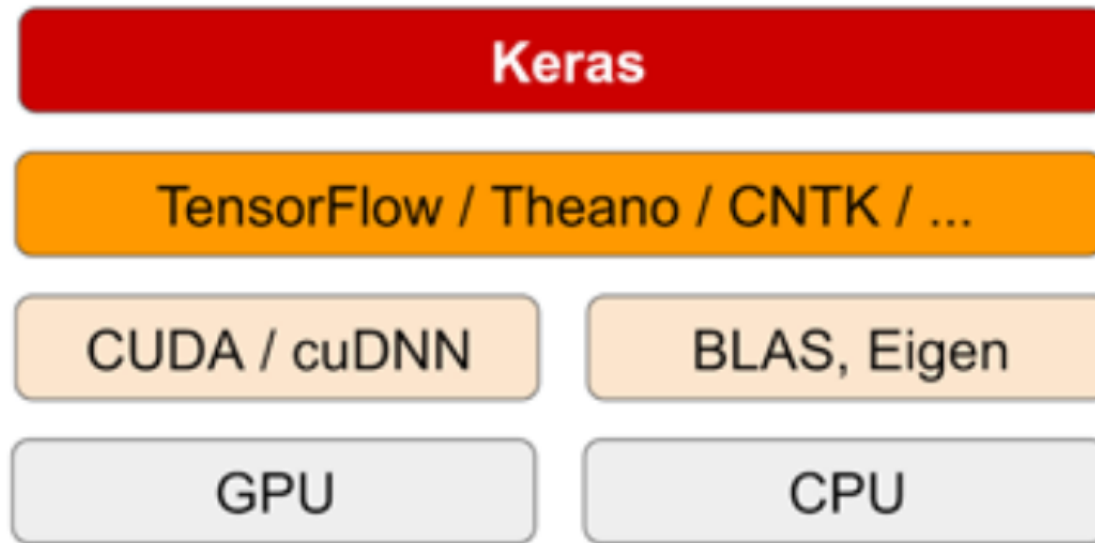
SDU Summer School

Deep Learning

Summer 2018

Introduction to KERAS


What is Keras?



- Keras is a high-level API providing easy to use elements for deep learning
- Can work with several backends
- Programs can easily deployed on CPUs, GPUs without changing the code

Who makes Keras? Contributors and backers

 **633** contributors



The Keras user experience

- Keras API is easy to understand.
 - Keras follows best practices for reducing cognitive load: it offers consistent & simple APIs, it minimizes the number of user actions required for common use cases, and it provides clear and actionable feedback upon user error.
- Easy to learn.
 - As a Keras user, you are more productive, allowing you to try more ideas than your competition, faster -- which in turn helps you win machine learning competitions.
- This ease of use does not come at the cost of reduced flexibility:
 - Keras integrates with lower-level deep learning languages (in particular TensorFlow), it enables you to implement anything you could have built in the base language. In particular, as `tf.keras`, the Keras API integrates seamlessly with your TensorFlow workflows.

Multi-Backend, Multi-Platform

- Develop in Python, R
 - On Unix, Windows, OSX
- Run the same code with...
 - TensorFlow
 - CNTK
 - Theano
 - MXNet
 - PlaidML
 - ??
- Run on CPU, NVIDIA GPU, AMD GPU, TPU...



How to use Keras: An introduction

Three API Styles

■ The Sequential Model

- Dead simple
- Only for single-input, single-output, sequential layer stacks
- Good for 70+% of use cases

■ The functional API

- Like playing with Lego bricks
- Multi-input, multi-output, arbitrary static graph topologies
- Good for 95% of use cases

■ Model subclassing

- Maximum flexibility
- Larger potential error surface

The Sequential API

```
import keras
from keras import layers

model = keras.Sequential()
model.add(layers.Dense(20, activation='relu', input_shape=(10,)))
model.add(layers.Dense(20, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

model.fit(x, y, epochs=10, batch_size=32)
```


The Functional API

```
import keras
from keras import layers

inputs = keras.Input(shape=(10,))
x = layers.Dense(20, activation='relu')(x)
x = layers.Dense(20, activation='relu')(x)
outputs = layers.Dense(10, activation='softmax')(x)

model = keras.Model(inputs, outputs)
model.fit(x, y, epochs=10, batch_size=32)
```

Model subclassing

```
import keras
from keras import layers

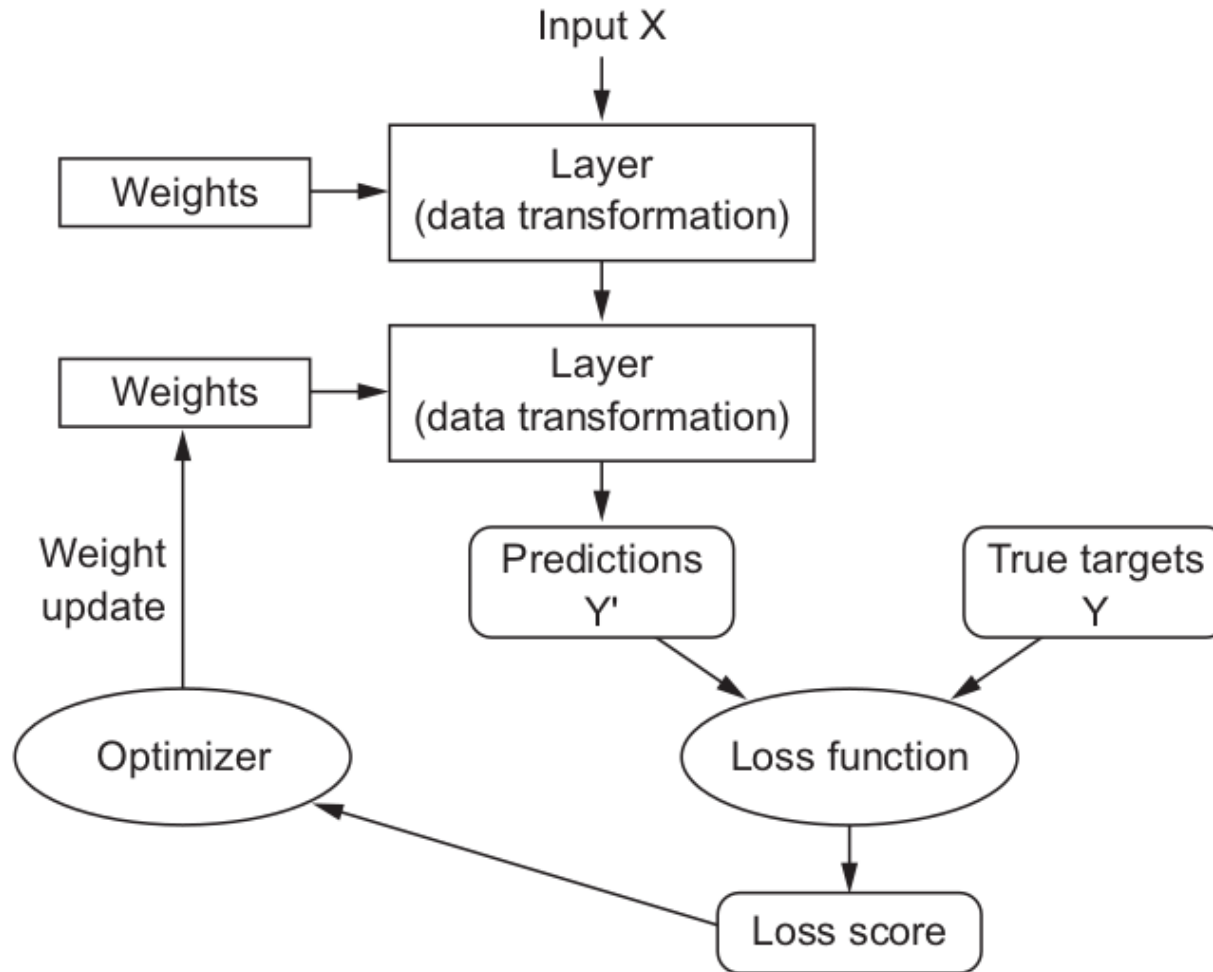
class MyModel(keras.Model):

    def __init__(self):
        super(MyModel, self).__init__()
        self.dense1 = layers.Dense(20, activation='relu')
        self.dense2 = layers.Dense(20, activation='relu')
        self.dense3 = layers.Dense(10, activation='softmax')

    def call(self, inputs):
        x = self.dense1(x)
        x = self.dense2(x)
        return self.dense3(x)

model = MyModel()
model.fit(x, y, epochs=10, batch_size=32)
```

Alrighty, let's put it together



Defining the Model

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Dense(32, activation='relu', input_shape=(784,)))
model.add(layers.Dense(10, activation='softmax'))
```

- The same model defined using the functional API:

```
input_tensor = layers.Input(shape=(784,))
x = layers.Dense(32, activation='relu')(input_tensor)
output_tensor = layers.Dense(10, activation='softmax')(x)

model = models.Model(inputs=input_tensor, outputs=output_tensor)
```

Options for Layers

- Core Layers
- Convolutional Layers
- Pooling Layers
- Locally-connected Layers
- Recurrent Layers
- Embedding Layers
- Merge Layers
- Normalization Layers
- Noise layers

Options for Layers

- The core layers perform the most basic operations
- They are enough to built FFN networks
- **Core Layers**
 - Dense Layers
 - Activation Layer
 - Dropout Layer
 - Reshape
 - Permute

Dense Layer

```
keras.layers.Dense(units,  
    activation=None, #Standard: use linear output  
    use_bias=True, #Add a bias vector  
    kernel_initializer='glorot_uniform', #How to initialize the  
                                     weights  
    bias_initializer='zeros', #How the biases  
    kernel_regularizer=None, #For example, apply L2 regularization  
    bias_regularizer=None, #For example, apply L2 regularization  
    activity_regularizer=None, #For example, apply L2 regularization  
    kernel_constraint=None, #For example, non-negative constraint  
    bias_constraint=None #For example, non-negative constraint  
)
```

Activation Function

```
model.add(Dense(64))  
model.add(Activation('tanh'))  
#This is equivalent to:  
model.add(Dense(64, activation='tanh'))
```

- Available Activations:
 - softmax
 - elu: (Exponential linear unit.)
 - x if $x > 0$ and $\alpha * (\exp(x) - 1)$ if $x < 0$.
 - selu: Scaled Exponential Linear Unit
 - softplus
 - $\log(\exp(x) + 1)$
 - relu
 - `relu(x, alpha=0.0, max_value=None)`
 - sigmoid
 - tanh

Compiling the Model

```
from keras import optimizers

model.compile(
    optimizer=optimizers.RMSprop(lr=0.001),
    loss='binary_crossentropy',
    metrics=['accuracy']
)
```

- Before training a model, you need to configure the learning process, which is done via the compile method, defining
 - **An optimizer.** This could be the string identifier of an existing optimizer (such as rmsprop or adagrad), or an instance of the Optimizer class
 - **A loss function.** This is the objective that the model will try to minimize. It can be the string identifier of an existing loss function, or it can be an objective function.
 - **A list of metrics.** A metric could be the string identifier of an existing metric or a custom metric function.

Examples: Compiling Models

```
# For a multi-class classification problem
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

```
# For a binary classification problem
model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

```
# For a mean squared error regression problem
model.compile(optimizer='rmsprop',
              loss='mse')
```

What does compile do?

- Compile defines the loss function, the optimizer and the metrics. That's all.
 - It has nothing to do with the weights and you can compile a model as many times as you want without causing any problem to pretrained weights.
 - You need a compiled model to train (because training uses the loss function and the optimizer). But it's not necessary to compile a model for predicting.
 - Do you need to use compile more than once? Only if:
 - You want to change one of these:
 - Loss function
 - Optimizer / Learning rate
 - Metrics
 - You loaded (or created) a model that is not compiled yet. Or your load/save method didn't consider the previous compilation.
- Consequences of compiling again:
 - If you compile a model again, you will lose the optimizer states.

Loss Functions

- `mean_squared_error`
- `mean_absolute_error`
- `mean_absolute_percentage_error`
- `mean_squared_logarithmic_error`
- `categorical_crossentropy`
- `sparse_categorical_crossentropy`
- `binary_crossentropy`
- ...

Metrics

- Can be any of the loss functions
- Some standard metrics like
 - F1
 - Precision
 - Recall
- have been removed in Keras 2.0
 - “These are all global metrics that were approximated batch-wise, which is more misleading than helpful.”

Train the Model

```
fit(x=None, y=None, # Input and desired outcome
    batch_size=None, # Number of samples per gradient update. If
                      # none, it defaults to 32
    epochs=1, # Number of runs over the complete x and y
    verbose=1, # Verbosity mode. 0 = silent, 1 = progress bar, 2 =
               # one line per epoch.
    callbacks=None, # List of functions to call during training
    validation_split=0.0, # Part of dataset set aside of test
    validation_data=None, # Validation dataset, tuple (x_val,y_val)
    shuffle=True, # shuffle the training data before each epoch
    class_weight=None, # Give some classes more/less weight
    sample_weight=None, # Give some samples more/less weight
    ...
)
```

In Context

```
model.compile(optimizer='rmsprop',  
              loss='binary_crossentropy',  
              metrics=['acc'])  
  
history = model.fit(partial_x_train,  
                    partial_y_train,  
                    epochs=20,  
                    batch_size=512,  
                    validation_data=(x_val, y_val))
```

The History Object

- Note that the call to `model.fit()` returns a History object. This object has a member `history`, which is a dictionary containing data about everything that happened during training.

```
>>> history_dict = history.history
>>> history_dict.keys()
[u'acc', u'loss', u'val_acc', u'val_loss']
```

- The dictionary contains four entries: one per metric that was being monitored during training and during validation.
- You can now plot these to get Information about your performance

Plotting the training and validation loss

```
import matplotlib.pyplot as plt


history_dict = history.history
loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']

epochs = range(1, len(loss_values) + 1)

# 'bo' is for blue dot, 'b' is for solid blue line
plt.plot(epochs, loss_values, 'bo', label='Training loss')
plt.plot(epochs, val_loss_values, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```

Plotting the training and validation loss



Example: Breast Cancer Revisited

- See demo

Larger Example with Cross-Validation

- Boston Housing Price data
 - You'll attempt to predict the median price of homes in a given Boston suburb in the mid-1970s
 - Given data points about the suburb at the time, such as the crime rate, the local property tax rate, and so on.
 - It has relatively few data points: only 506, split between 404 training samples and 102 test samples.
 - Each feature in the input data (for example, the crime rate) has a different scale.
- We have to solve a regression problem
- See demo

Practical Recommendations

- The fewer data, you should train smaller and shallower networks in order to prevent overfitting
- Preprocessing
 - Take small values - Typically, most values should be in the 0–1 range.
 - Be homogenous- That is, all features should take values in roughly the same range.

Regularization

- Adding weight regularization

```
from keras import regularizers

model = models.Sequential()
model.add(layers.Dense(16,
    kernel_regularizer=regularizers.l2(0.001),
    activation='relu', input_shape=(10000,)))

model.add(layers.Dense(1, activation='sigmoid'))
```

- `l2(0.001)` means every coefficient in the weight matrix of the layer will add $0.001 * \text{weight_coefficient_value}$ to the total loss of the network.
- Note that because this penalty is only added at training time, the loss for this network will be much higher at training than at test time.

Dropout

```
model = models.Sequential()

model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dropout(0.5))

model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dropout(0.5))

model.add(layers.Dense(1, activation='sigmoid'))
```

Load and Save models

- You save a Keras model into a single HDF5 file which will contain:
 - the architecture of the model, allowing to re-create the model
 - the weights of the model
 - the training configuration (loss, optimizer)
 - the state of the optimizer, allowing to resume training exactly where you left off.

```
from keras.models import load_model

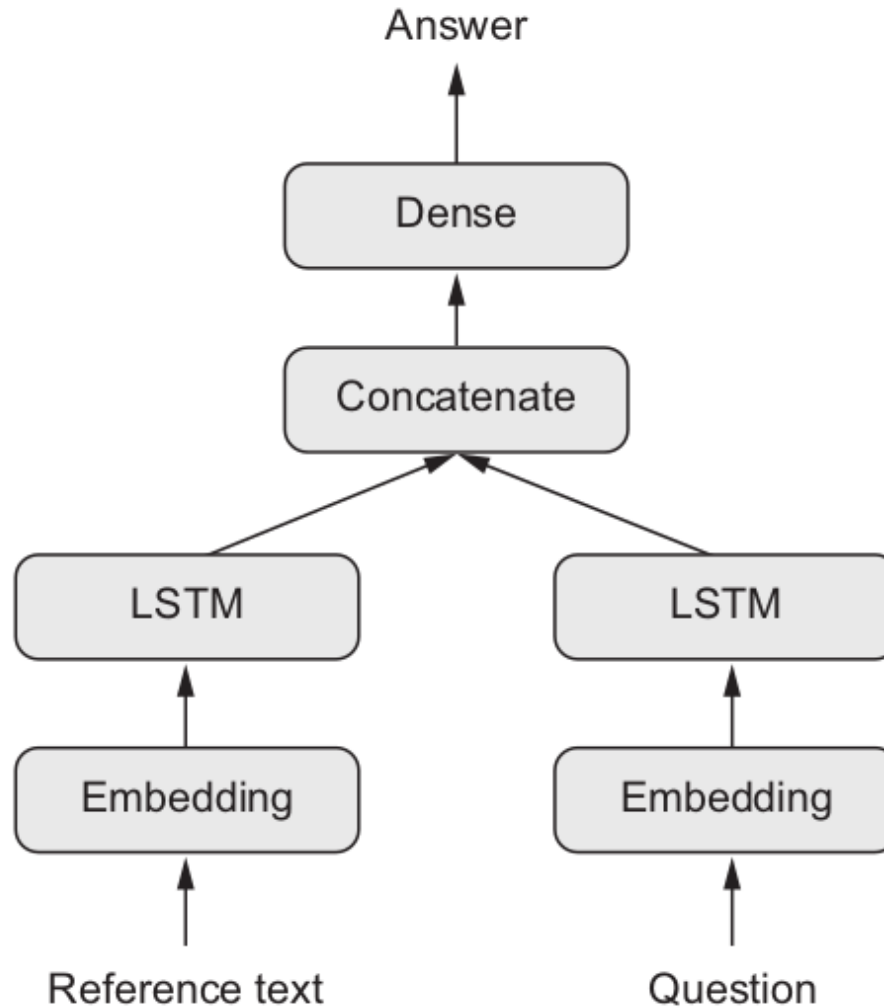
model.save('my_model.h5') # creates a HDF5 file 'my_model.h5'
del model # deletes the existing model

# returns a compiled model
# identical to the previous one
model = load_model('my_model.h5')
```


Implementing your own Callback

- Callbacks are implemented by sub-classing the class `keras.callbacks.Callback`.
- You can implement the following functions:
 - `on_epoch_begin`
 - `on_epoch_end`
 - `on_batch_begin`
 - `on_batch_end`
 - `on_train_begin`
 - `on_train_end`

Multi Input Models



Multi Input Model using the functional API

```
text_input = Input(shape=(None,), dtype='int32', name='text')
embedded_text = layers.Embedding(64, text_vocabulary_size)(text_input)
encoded_text = layers.LSTM(32)(embedded_text)

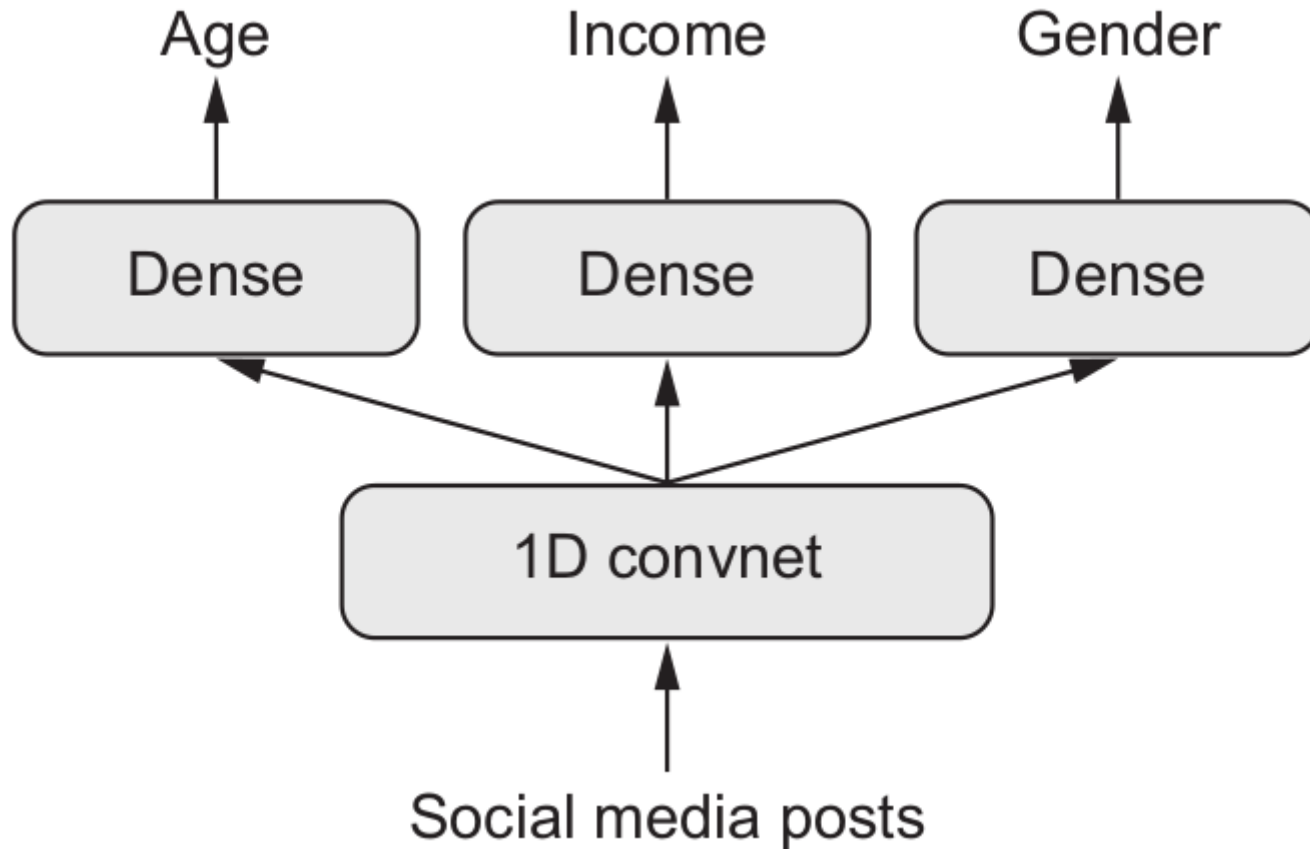
question_input = Input(shape=(None,), dtype='int32', name='question')
embedded_question = layers.Embedding(32,
    question_vocabulary_size)(question_input)
encoded_question = layers.LSTM(16)(embedded_question)

concatenated = layers.concatenate([encoded_text, encoded_question],
    axis=-1)

answer = layers.Dense(answer_vocabulary_size,
    activation='softmax')(concatenated)

model = Model([text_input, question_input], answer)
model.compile(optimizer='rmsprop', loss='categorical_crossentropy',
    metrics=['acc'])
```

Multi-Output Models



|

Multi-Output Models

```
posts_input = Input(shape=(None,), dtype='int32', name='posts')
embedded_posts = layers.Embedding(256, vocabulary_size)(posts_input)
x = layers.Conv1D(128, 5, activation='relu')(embedded_posts)
x = layers.MaxPooling1D(5)(x)
... # Construct the network how you like it
x = layers.Dense(128, activation='relu')(x)

age_prediction = layers.Dense(1, name='age')(x)
income_prediction = layers.Dense(num_income_groups, activation='softmax',
                                name='income')(x)
gender_prediction = layers.Dense(1, activation='sigmoid', name='gender')(x)

model = Model(posts_input,[age_prediction, income_prediction,
                             gender_prediction])
```

Multi-Output Models

```
model.compile(optimizer='rmsprop',  
              loss=['mse', 'categorical_crossentropy', 'binary_crossentropy'])
```

```
model.compile(optimizer='rmsprop',  
              loss={'age': 'mse',  
                    'income': 'categorical_crossentropy',  
                    'gender': 'binary_crossentropy'})
```

- Fitting, etc, of the model remains the same as with a normal network.

Callbacks

- When training a model, there are many things you cannot predict
- Sometime it would be helpful to intervene when something goes wrong
- Keras provides callbacks:
 - Model checkpointing: Saving the current weights of the model at different points during training.
 - Early stopping: Interrupting training when the validation loss is no longer improving (and of course, saving the best model obtained during training).
 - Dynamically adjusting the value of certain parameters during training: Such as the learning rate of the optimizer.
 - Logging training and validation metrics during training, or visualizing the representations learned by the model as they're updated: The Keras progress bar is a callback!

Early Stopping

```
import keras

callbacks_list = [
    keras.callbacks.EarlyStopping(
        monitor='acc',
        patience=1
    ),
    keras.callbacks.ModelCheckpoint(
        filepath='my_model.h5',
        monitor='val_loss',
        save_best_only=True
    )
]

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['acc'])

model.fit(x, y, epochs=10, batch_size=32, callbacks=callbacks_list,
          validation_data=(x_val, y_val))
```


Deep Learning Workflow

1. Defining the problem and assembling a dataset
2. Choosing a measure of success
3. Decide on the evaluation protocol
4. Preparing your data
5. Define a model better than base-line
6. Scaling up: Make the model overfit
7. Regularizing your model and tuning your hyperparameters
8. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset

- What will your input data be?
- What type of problem are you facing? Classification? Regression? ...

2. Choosing a measure of success

3. Decide on the evaluation protocol

4. Preparing your data

5. Define a model better than base-line

6. Scaling up: Make the model overfit

7. Regularizing your model and tuning your hyperparameters

8. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
- 2. Choosing a measure of success**
 - How do you measure if the model is successful
 - Not to be confused with the loss function which is often only a surrogate for what you actually want achieve
3. Decide on the evaluation protocol
4. Preparing your data
5. Define a model better than base-line
6. Scaling up: Make the model overfit
7. Regularizing your model and tuning your hyperparameters
8. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
2. Choosing a measure of success
3. **Decide on the evaluation protocol**
 - hold-out validation
 - Cross-validation
4. Preparing your data
5. Define a model better than base-line
6. Scaling up: Make the model overfit
7. Regularizing your model and tuning your hyperparameters
8. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
2. Choosing a measure of success
3. Decide on the evaluation protocol
4. **Preparing your data**
 - Clean data
 - Normalize data
5. Define a model better than base-line
6. Scaling up: Make the model overfit
7. Regularizing your model and tuning your hyperparameters
8. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
2. Choosing a measure of success
3. Decide on the evaluation protocol
4. Preparing your data
5. **Define a model better than base-line**
 - Last-layer activation
 - Loss function
 - Optimization Algorithm
6. Scaling up: Make the model overfit
7. Regularizing your model and tuning your hyperparameters
8. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
2. Choosing a measure of success
3. Decide on the evaluation protocol
4. Preparing your data
5. Define a model better than base-line
6. **Scaling up: Make the model overfit**
 - Add layers.
 - Make the layers bigger.
 - Train for more epochs.
7. Regularizing your model and tuning your hyperparameters
8. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
2. Choosing a measure of success
3. Decide on the evaluation protocol
4. Preparing your data
5. Define a model better than base-line
6. Scaling up: Make the model overfit
- 7. Regularizing your model and tuning your hyperparameters**
 - This will take the most time
 - Add dropout.
 - Try different architectures: add or remove layers.
 - Add L1 and/or L2 regularization
8. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
2. Choosing a measure of success
3. Decide on the evaluation protocol
4. Preparing your data
5. Define a model better than base-line
6. Scaling up: Make the model overfit
7. Regularizing your model and tuning your hyperparameters
8. **Finalize your final model**
 - Save and distribute the model